

Carterfone

The switchboard for personal AI agents.

A neutral, security-first connectivity layer letting AI agents across vendors interoperate — under permissions, approvals, identity proofs, and audit.

Shanto Mathew, Founder · SprintSite LLC · Aubrey, Texas
shanto12@gmail.com · shantomathew.com · github.com/shanto12
github.com/shanto12/carterfone

Agents can't talk to each other safely.

- AI assistants are multiplying across vendors — ChatGPT, Claude, open models, custom agents.
- They cannot interoperate: every connection is a brittle, custom, one-off integration.
- Worse, agents today share raw credentials and API keys with no permissions, no approvals, no audit trail.
- As agents start touching money, data, and decisions, that is a security incident waiting to happen.

One trusted switchboard for every agent.

- Carterfone is a neutral connectivity layer any agent can join, regardless of vendor.
- Every cross-agent call runs under explicit permissions and approvals.
- Agents carry verifiable identity proofs — you always know who acted.
- Every action lands in an audit log. Security-first by architecture, not by policy document.

Built on open protocols. Open repo.

- Speaks A2A and MCP — the emerging open standards for agent communication.
- TypeScript implementation. Public repo: github.com/shanto12/carterfone
- Architecture and specification phase complete; implementation actively underway.
- Designed as infrastructure: the trust layer other agent products build on.

Three steps. Every call.

1. Agents register

Each agent gets a verifiable identity proof before it can touch the switchboard.

2. Calls are authorized

Every cross-agent request passes permissions checks and required approvals.

3. Actions are audited

Each action is written to an immutable audit log, tied to the agent's identity.

Agents are moving into real workflows.

- Agents are graduating from chat demos to handling data, money, and operational decisions.
- Interoperation without a trust layer does not scale — every pair of agents needs its own custom plumbing.
- Security, identity, and audit are the missing infrastructure for the agent economy.

Market sizing in progress — no TAM claims made in this deck.

Building now. Pre-launch.

- Status: architecture and spec complete; implementation underway (repo created October 2026, actively developed).
- No users yet. No revenue. No outside funding raised.
- Current focus: shipping the core switchboard and the permissions/approvals engine.

Everything in this deck reflects the actual state of the project — nothing inflated.

In exploration.

- No pricing decided. The focus now is shipping the trust layer, not monetizing it.
- Directions under consideration (not decided): usage-based pricing for cross-agent calls.
- Directions under consideration (not decided): an enterprise tier for compliance, audit, and SSO.

Business model will be set after private-beta feedback.

Solo founder, shipping.

- Shanto Mathew — Founder. AI engineer, Forward Deployed / AI Engineer background.
- Experience across LLMs, voice AI, agent frameworks including LangGraph, and RAG systems.
- Based in Aubrey, Texas. Entity: SprintSite LLC (Texas LLC).
- Hiring: TBD — technical hires after pre-seed.

Pre-seed funding to ship the switchboard.

- Use of funds: complete the implementation, harden the security model, launch a private beta.
- Raising \$500K pre-seed. Terms: TBD.
- Contact: shanto12@gmail.com · shantomathew.com · github.com/shanto12